



Aalto University  
School of Electrical  
Engineering

# ELEC-E8125 Reinforcement Learning

## Solving discrete MDPs


Joni Pajarinen

12.9.2023

# Previous lecture: find shortest path exercise

- Use backward value iteration for

Initial start state  $\longrightarrow s_I = S$   
 $S_G = \{B\}$



Final cost  

$$l_F(s) = \begin{cases} 0, & s \in S_G \\ \infty, & s \notin S_G \end{cases}$$
  
 Goal set

Reminder:

$$G^*(s) = \min_a \{l(s, a) + G^*(f(s, a))\}$$


# Today

- Markov decision processes (MDPs)

# Learning goals

- Understand MDPs and related concepts
- Understand value functions
- Be able to implement value iteration for determining an optimal policy (in exercise #2, you will get to do this for real)

# Markov decision process (MDP)



## MDP

Environment observable

$$z = s$$

Defined by dynamics

$$P(s_{t+1} | s_t, a_t)$$

And reward function

$$r_t = r(s_t, a_t)$$

Solution, for example

$$a_{1, \dots, T}^* = \arg \max_{a_1, \dots, a_T} \sum_{t=1}^T r_t$$

Represented as policy

$$a = \pi(s)$$

# Markov property



Andrey Markov

- “Future is independent of past given the present”
- State sequence  $S$  is Markov iff  $\longleftrightarrow$  “if and only if”


$$P(S_{t+1}|S_t) = P(S_{t+1}|S_1, \dots, S_t)$$

- State captures all history
- Once state is known, history may be thrown away

# Markov process

No “decision” here!

- Markov process is a memoryless random process that generates a state sequence  $S$  with the Markov property
- Defined as  $(S, T)$ 
  - $S$ : set of states
  - $T: S \times S \rightarrow [0, 1]$  state transition function
    - $T_t(s, s') = P(s_{t+1} = s' | s_t = s)$
    - $P$  can be represented as a transition probability matrix
- State sequences called *episodes*




Still no “decision”!

# Markov reward process

- Markov reward process =  
Markov process with rewards
- Defined by  $(S, T, r, \gamma)$ 
  - $S, T$ : as above
  - $r: S \rightarrow \mathcal{R}$  reward function
  - $\gamma \in [0, 1]$  discount factor
- Accumulated rewards in finite ( $H$  steps) or infinite horizon

$$\sum_{t=0}^H \gamma^t r_t \quad \sum_{t=0}^{\infty} \gamma^t r_t$$

- *Return*  $G$ : accumulated rewards from time  $t$



$$G_t = \sum_{k=0}^H \gamma^k r_{t+k}$$

Why discount?

Return of (A,B,C),  $\gamma = 0.9$ ?



# State value function for Markov reward processes


- State value function  $V(s)$  is expected cumulative reward starting from state  $s$

$$V(s) = E[G_t | s_t = s]$$

- Value function can be defined by the Bellman equation

$$V(s) = E[G_t | s_t = s]$$

$$V(s) = E[r_t + \gamma V(s_{t+1}) | s_t = s]$$



What is the value function for  $\gamma = 0$ ?

What is the value function for  $\gamma = 0.5$  after a single Bellman update when starting with zero values?

# Markov decision process (MDP)

- Markov decision process defined by  $(S, A, T, R, \gamma)$ 
  - $S, \gamma$ : as above
  - $A$ : set of actions (inputs)
  - $T: S \times A \times S \rightarrow [0, 1]$ 

$$T_t(s, a, s') = P(s_{t+1} = s' | s_t = s, a_t = a)$$
  - $R: S \times A \rightarrow \mathcal{R}$  reward function
 
$$r_t(s, a) = r(s_t = s, a_t = a)$$
- Goal: Find policy  $\pi(s)$  that maximizes expected cumulative reward

Grid world

			+1
			-1

Agent tries to move forward:

$P(\text{success}) = 0.8$

$P(\text{left}) = 0.1$

$P(\text{right}) = 0.1$

0.1	
→	0.8
0.1	

	0.8	
0.1	↑	0.1

# Policy

- Deterministic policy  $\pi(S): S \rightarrow A$  is a mapping from states to actions
- Stochastic policy  $\pi(a|s): S, A \rightarrow [0, 1]$  is a distribution over actions given states
- Optimal policy  $\pi^*(s)$  is a policy that is better or equal than any other policy (in terms of cumulative rewards)
  - There always exists a deterministic optimal policy for an MDP

Grid world

			+1
			-1

Agent tries to move forward:

$P(\text{success}) = 0.8$

$P(\text{left}) = 0.1$

$P(\text{right}) = 0.1$

0.1	
→ 0.8	
0.1	

	0.8	
0.1	↑	0.1

# MDP value function

- *State-value function* of an MDP is the expected return starting from state  $s$  and following policy  $\pi$

$$V_{\pi}(s) = E_{\pi}[G_t | s_t = s]$$

- Can be decomposed into immediate and future components using Bellman expectation equation

$$V_{\pi}(s) = E_{\pi}[r_t + \gamma V_{\pi}(s_{t+1}) | s_t = s]$$

$$V_{\pi}(s) = r(s, \pi(s)) + \gamma \sum_{s'} T(s, \pi(s), s') V_{\pi}(s')$$

			+1
			-1

→	→	→	X
↑		↑	↑
↑	→	↑	←

# Action-value function

- *Action-value function*  $Q$  is expected return starting from state  $s$ , taking action  $a$ , and then following policy  $\pi$

$$Q_{\pi}(s, a) = E_{\pi}[G_t | s_t = s, a_t = a]$$

- Using Bellman expectation equation

$$Q_{\pi}(s, a) = E_{\pi}[r_t + \gamma Q_{\pi}(s_{t+1}, a_{t+1} | s_t = s, a_t = a)]$$

$$Q_{\pi}(s, a) = r(s, a) + \gamma \sum_{s'} T(s, a, s') Q_{\pi}(s', \pi(s'))$$

			+1
			-1

→	→	→	X
↑		↑	↑
↑	→	↑	←

# Optimal value function

- Optimal state-value function is maximum value function over all policies

$$V^*(s) = \max_{\pi} V_{\pi}(s)$$

- Optimal action-value function is maximum action-value function over all policies

$$Q^*(s, a) = \max_{\pi} Q_{\pi}(s, a)$$

- All optimal policies achieve optimal state- and action-value functions

# Optimal policy vs optimal value function

- Optimal policy for optimal action-value function

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$

- Optimal action for optimal state-value function

$$\pi^*(s) = \arg \max_a E_{s'}[r(s, a) + \gamma V^*(s')]$$

$$\pi^*(s) = \arg \max_a \left( r(s, a) + \gamma \sum_{s'} T(s, a, s') V^*(s') \right)$$

# Value iteration

Do you notice that this is an expectation?

- Starting from  $V_0^*(s) = 0 \quad \forall s$   
iterate

$$V_{i+1}^*(s) = \max_a \left( r(s, a) + \gamma \sum_{s'} T(s, a, s') V_i^*(s') \right)$$

until convergence

- Value iteration converges to  $V^*(s)$

Compare to

$$G^*(s) = \min_a \{ l(s, a) + G^*(f(s, a)) \}$$


from last week!



# Iterative policy evaluation

- Problem: Evaluate value of policy  $\pi$
- Solution: Iterate Bellman expectation back-ups
- $V_1 \rightarrow V_2 \rightarrow \dots \rightarrow V_\pi$
- Using synchronous back-ups:
  - For all states  $s$
  - Update  $V_{k+1}(s)$  from  $V_k(s')$
  - Repeat

From slide 12



$$V_{k+1}(s) = r(s, \pi(s)) + \gamma \sum_{s'} T(s, \pi(s), s') V_k(s')$$

$$V_{k+1}(s) = \sum_a \pi(a|s) \left( r(s, a) + \gamma \sum_{s'} T(s, a, s') V_k(s') \right)$$

V

Greedy policy

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

	↕	↕	↕
↕	↕	↕	↕
↕	↕	↕	↕
↕	↕	↕	

	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

	←	↕	↕
↑	↕	↕	↕
↕	↕	↕	↓
↕	↕	→	

$r = -1$  for all actions

$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

	←	←	↕
↑	↖	↕	↓
↑	↕	↘	↓
↕	→	→	

# Policy improvement and policy iteration

- Given a policy  $\pi$ , it can be improved by
  - Evaluating  $V_\pi$
  - Forming a new policy by acting greedily with respect to  $V_\pi$
- This always improves the policy
- Iterating multiple times called *policy* iteration
  - Converges to optimal policy

# Computational limits – Value iteration

- Complexity  $O(|A||S|^2)$  per iteration
- Effective up to medium size problems (millions of states)
- Complexity when applied to action-value function  $O(|A|^2|S|^2)$  per iteration

# Summary

- Markov decision processes represent environments with uncertain dynamics
- Deterministic optimal policies can be found using state-value or action-value functions
- Dynamic programming is used in value iteration and policy iteration algorithms

# Next week: From MDPs to RL

- Readings
  - Sutton & Barto Ch. 5-5.4, 5.6, 6-6.5