



Aalto University  
School of Electrical  
Engineering

# ELEC-E8125 Reinforcement Learning

## Reinforcement learning in discrete domains

Joni Pajarinen

19.9.2023


# Today

- Reinforcement learning
- Policy evaluation vs control problems
- Monte-Carlo and Temporal difference

# Learning goals

- Understand basic concepts of RL
- Understand Monte-Carlo (MC) and temporal difference (TD) approaches for policy evaluation and control
- Be able to implement MC and TD

# Reinforcement learning



## RL

Markov decision process (MDP) with unknown Markovian dynamics

$$P(s_{t+1}|s_t, a_t)$$

Unknown reward function

$$r_t = r(s_t, a_t)$$

Solution similar, e.g.

$$a_{1,\dots,T}^* = \max_{a_1,\dots,a_T} \sum_{t=1}^T r_t$$

Learning must explore policies

# Reinforcement learning

- MDP with unknown dynamics ( $P$ ) and reward function ( $r$ )
- Model based RL: Estimate MDP, apply MDP methods
  - Estimate MDP dynamics and reward functions from data
- Can we do without  $P$  and  $r$  ?
  - Can we evaluate a policy (construct value function) if we have multiple episodes (in episodic tasks) available?

# Monte-Carlo policy evaluation

- Complete episodes give us samples of return  $G$
- Learn value of particular policy from episodes under that policy:

$$V_{\pi}(s) = E_{\pi}[G_t | s_t = s] \quad G_t = \sum_{k=0}^H \gamma^k r_{t+k}$$

- Estimate value as empirical mean return

– For each visited state  $s$  in an episode,

$$N(s) = N(s) + 1 \quad S(s) = S(s) + G_t \quad V(s) \approx S(s) / N(s)$$

- When number of episodes approaches infinity,

$V(s)$  converges:  $V(s) \rightarrow V_{\pi}(s)$



Empirical mean approaches true mean.  
Can we do without episodes?

# Temporal difference (TD) – learning without episodes

- For each state transition, update a guess towards a guess:

$$V(s_t) = V(s_t) + \alpha (r_t + \gamma V(s_{t+1}) - V(s_t))$$

- Approach called TD(0)

Estimated return



- Compare to MC

$$V(s_t) = V(s_t) + \alpha (G_t - V(s_t))$$

True return



# Batch learning

- For a limited number of trials available:
  - Sample episode  $k$
  - Apply MC or TD(0) to episode  $k$ .

A, 0, B, 0

B, 1

B, 1


B, 1

B, 1

B, 1

B, 1

B, 0



What is  $V(A)$ ?



# MC vs TD


- MC
  - Needs full episodes. Only works in episodic environments
  - High variance, zero bias → good but slow convergence
  - Does not exploit Markov property → often better in non-Markov environments
- TD (esp. TD(0))
  - Can learn from incomplete episodes and online after each time step
  - Works in continuing non-episodic environments
  - Low variance, some bias → often more efficient than MC, discrete state TD(0) converges, more sensitive to initial value
  - Exploits Markov property → often more efficient in Markov environments

# $\lambda$ -return

- Combine returns in different horizons

$$G_t^\lambda = (1 - \lambda) \sum_{k=0}^{\infty} \lambda^k G_t^{(k)}$$

k-step return:  $G_t^{(k)} = \sum_{i=0}^{k-1} \gamma^i r_{t+i} + \gamma^k V(s_{t+k})$



- State value function update becomes

$$V(s_t) = V(s_t) + \alpha (G_t^\lambda - V(s_t))$$

- This is called TD( $\lambda$ ) based on the *forward view*

# Causes and effects – eligibility traces

- Which state is the “cause” of a reward?
- Frequency heuristic: most frequent states likely
- Recency heuristic: most recent states likely

- *Eligibility trace* for a state combines these:

$$E_t(s) = \gamma \lambda E_{t-1}(s) + \mathbf{1}(s_t = s) \quad \leftarrow \text{“How often a particular state was visited recently on average?”}$$

Not exactly, also considers discount.

# Backward-TD( $\lambda$ )

- Extend TD time horizon with decay ( $\lambda$ )

- After episode, update

$$V(s) = V(s) + \alpha E_t(s) (r_t + \gamma V(s_{t+1}) - V(s_t))$$

- TD(1) equal to MC

$$E_t(s) = \gamma \lambda E_{t-1}(s) + \mathbf{1}(s_t = s)$$

- Eligibility traces way to implement *backward* TD( $\lambda$ ),  
*forward* TD( $\lambda$ ) requires episodes

# Control / decision making?

- So far we only found out how to estimate value functions for a particular policy
- Can we use this to optimize a policy?

# Policy improvement and policy iteration

- Given a policy  $\pi$ , it can be improved by
  - Evaluating its value function
  - Forming a new policy by acting greedily with respect to the value function
- This always improves the policy
- Iterating multiple times called *policy* iteration
  - Converges to optimal policy

# Monte-Carlo Policy iteration

- Can we choose action using value function  $V(s)$  ?
- Greedy policy improvement using action-value function  $Q(s,a)$  does not require a model:

$$\pi'(s) = \arg \max_a Q(s, a)$$

- Estimate  $Q(s,a)$  using MC (empirical mean = “calculate average”)

Note: calculate frequencies for all state-action pairs.

# Ensuring exploration

Note: in exercise 3 task 1, we will use epsilon greedy exploration

- Simple approach:  $\epsilon$ -greedy exploration:
  - Explore: Choose action at random with probability  $\epsilon$
  - Exploit: Be greedy with probability  $1 - \epsilon$

$$\pi(a|s) = \begin{cases} 1 - \epsilon & \text{if } a = \arg \max_a Q(s, a') \\ \epsilon / (N_{\text{actions}} - 1) & \text{for any other action} \end{cases}$$

- How to converge to optimal policy?
    - Idea: reduce  $\epsilon$  over time.
    - For example, for  $k$ :th episode  $\epsilon = \frac{b}{b + k}$
- Number of different actions
- “Greedy in the Limit with Infinite Exploration” (GLIE)
- constant



# SARSA


- Idea: Apply TD to  $Q(S,A)$ 
  - With  $\varepsilon$ -greedy policy improvement
  - Update each time step

$$Q(s, a) = Q(s, a) + \alpha (r + \gamma Q(s', a') - Q(s, a))$$

Compare with

$$V(s_t) = V(s_t) + \alpha (r_t + \gamma V(s_{t+1}) - V(s_t))$$

- SARSA converges under
  - GLIE policy (greedy in the limit of infinite exploration),
  - $\sum_{t=0}^{\infty} \alpha_t = \infty \quad \sum_{t=0}^{\infty} \alpha_t^2 < \infty$



# SARSA( $\lambda$ )

- Instead of TD(0) update in SARSA, use TD( $\lambda$ ) update
- Backward SARSA( $\lambda$ )

$$E_t(s, a) = \gamma \lambda E_{t-1}(s, a) + \mathbf{1}(s_t = s, a_t = a)$$


$$Q(s, a) = Q(s, a) + \alpha E_t(s, a) (r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$$

Compare to


$$E_t(s) = \gamma \lambda E_{t-1}(s) + \mathbf{1}(s_t = s)$$

$$V(s) = V(s) + \alpha E_t(s) (r_t + \gamma V(s_{t+1}) - V(s_t))$$


Path taken



Action values increased  
by one-step Sarsa



Action values increased  
by Sarsa( $\lambda$ ) with  $\lambda=0.9$



# On-policy vs off-policy learning

- *On-policy learning* (methods so far)
  - Use a policy while learning how to optimize it
  - “Learn on the job”
- *Off-policy learning*
  - Use another policy while learning about optimal policy
  - Can learn from observation of other agents
  - Can learn about optimal policy when using exploration policy

Note: in exercise 3 task 1,  
we will implement Q-learning

# Q-learning

- Use  $\varepsilon$ -greedy *behavior policy* to choose actions
- *Target policy* is greedy with respect to  $Q$

$$\pi(s) = \arg \max_a Q(s, a)$$

- Update target policy greedily:

$$Q(s, a) = Q(s, a) + \alpha (r + \gamma \max_{\hat{a}} Q(s', \hat{a}) - Q(s, a))$$

- $Q$  converges to  $Q^*$

Assume we take greedy action at next step.



# Summary

- In reinforcement learning, dynamics and reward function of the MDP are in general unknown
- MC (Monte-Carlo) approaches sample returns from full episodes
- TD (temporal difference) approaches sample estimated returns (biased)
- Returns can be used to update a policy or value function

# Next: Extending state spaces

- What to do if
  - discrete state space is too large?
  - state space is continuous?
- Readings
  - Sutton & Barto, ch. 9-9.3, 10-10.1