



Aalto University
School of Electrical
Engineering

ELEC-E8125 Reinforcement Learning

Interleaved learning and planning in model-based RL

Joni Pajarinen

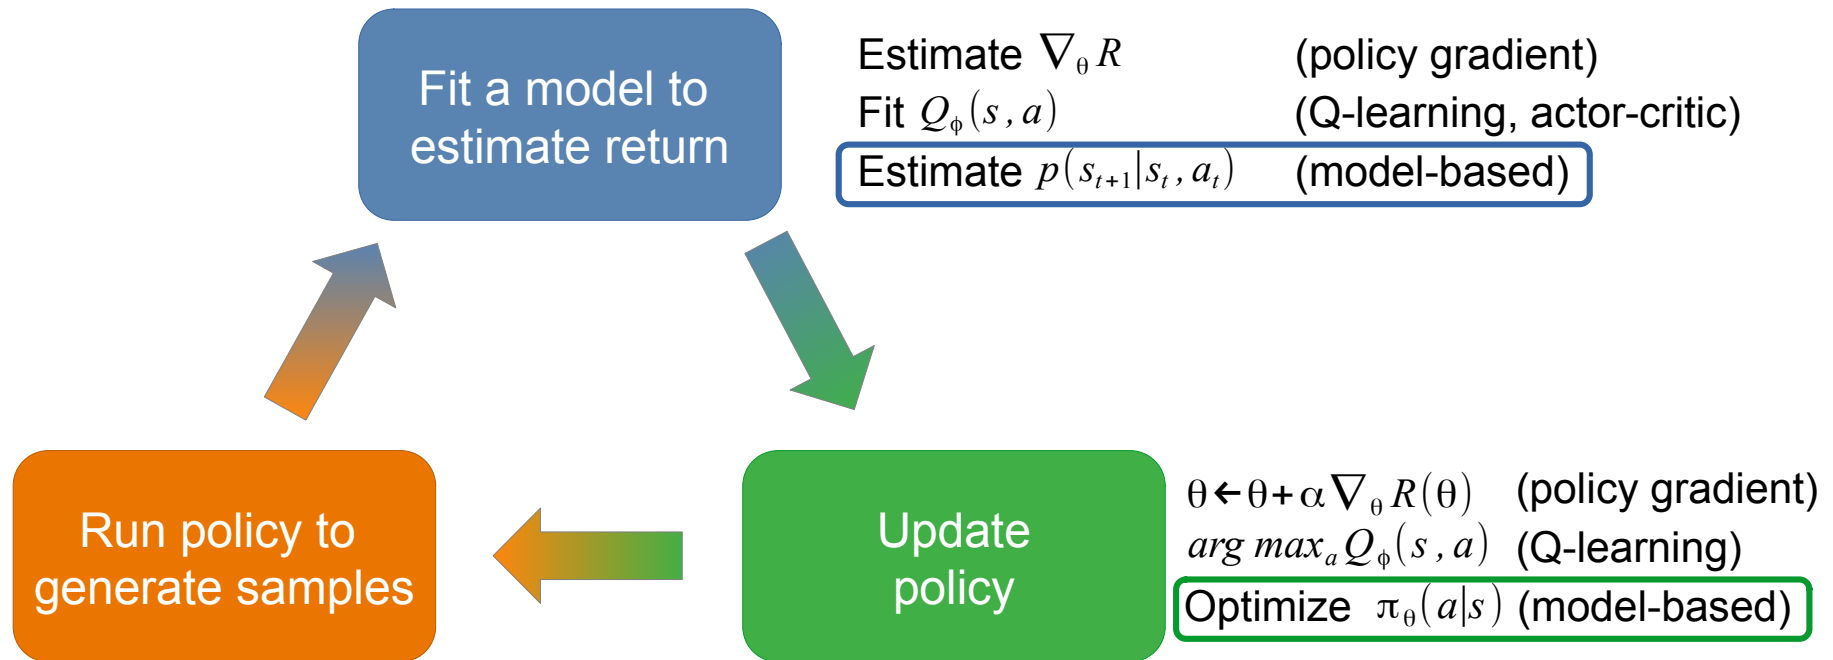
31.10.2023

Learning goals

- Understand how learning and planning are used together in model-based reinforcement learning

Anatomy of reinforcement learning

Model-based



Motivation (partial recap)



- Reinforcement learning has limited sample efficiency
- Learned policies are task(reward-function)-specific, learned policies cannot be directly reused
- Learned dynamics model is reusable and can be used to reason about potential futures
- Sometimes we know the model, e.g. in games!



Model definition and types

- Dynamics model $s_{t+1} = f(s_t, a_t)$ or $f(s_{t+1} | s_t, a_t)$
- Reward model $r_t = r(s_t, a_t)$ or $r(r_t | s_t, a_t)$
- Models are usually learned
 - Parametric regression (e.g. neural net) common
- May be also known (e.g. games, simulators)
 - Even physics based models need to be often calibrated
- Also other possibilities (active research area)
 - Latent variable models, graph neural networks, non-parametric regression models such as Gaussian processes, ...

Which model to use?



$$Y_i = \beta_0 + \beta_1 \phi_1(X_{i1}) + \dots + \beta_p \phi_p(X_{ip})$$

Gaussian process (GP)

- Data-efficient
- Slow with big datasets
- May be too smooth for non-smooth dynamics

Neural networks (NNs)

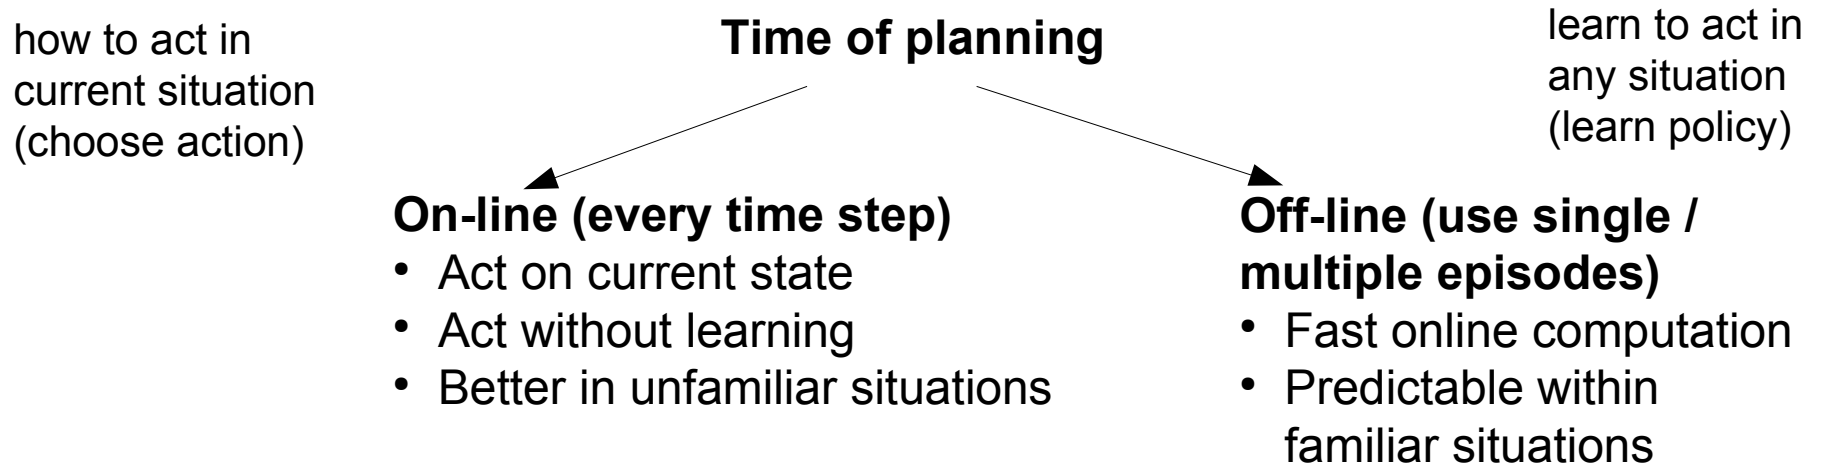
- Expressive
- Unpredictable with sparse data (overfit)
 - NN ensembles estimate uncertainty

Linear models


- May be used locally
- Do not overfit

Domain specific parametric models (e.g. physics parameters) can also be used
→ Traditional control engineering approach of model identification + control

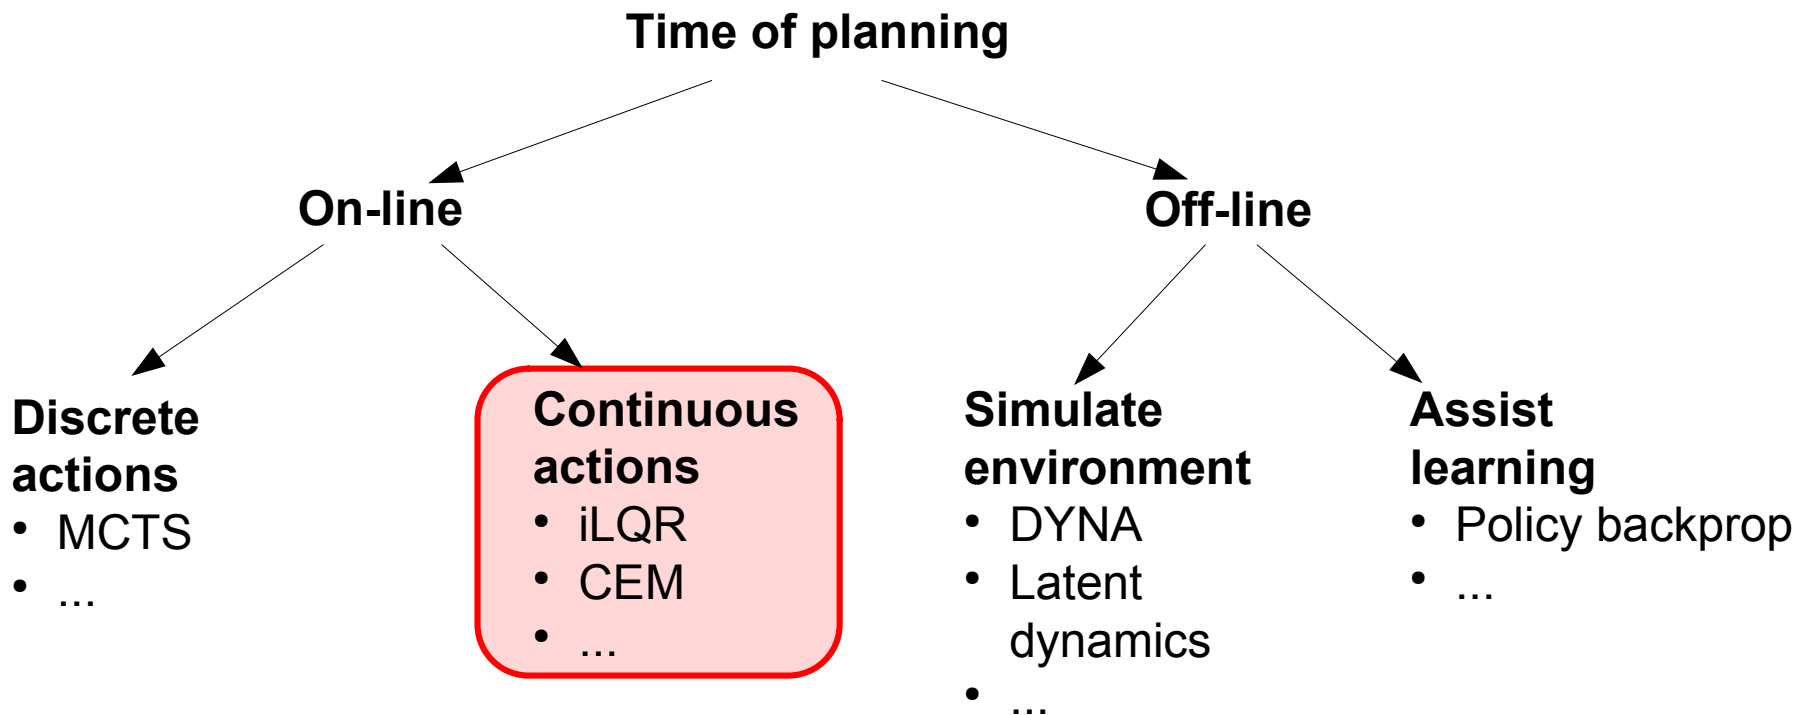
Spectrum of model-based RL



Spectrum of model-based RL



Spectrum of model-based RL



Continuous on-line planning: iLQR + learned model

Input: base policy π_0

Run base policy to collect data $D \leftarrow \{(s, a, s')_i\}$

Repeat

Fit dynamics model $f(s, a)$ to minimize $\sum_i \|f(s_i, a_i) - s'_i\|^2$

Use model to plan (e.g. iLQR, CEM) actions

Execute first planned action, observe resulting state s'

Update dataset $D \leftarrow D \cup \{(s, a, s')\}$

Continuous on-line planning: iLQR + learned model

Input: base policy π_0

Run base policy to collect data $D \leftarrow \{(s, a, s')_i\}$

Repeat

Fit dynamics model $f(s, a)$ to minimize $\sum_i \|f(s_i, a_i) - s'_i\|^2$

Use model to plan (e.g. iLQR, CEM) actions

Execute first planned action, observe resulting state s'

Update dataset $D \leftarrow D \cup \{(s, a, s')\}$

- Sample efficient
- Computationally expensive for two reasons
 - Dynamics fitting costly \rightarrow model may be fitted only periodically (every n steps)
 - Planning costly for long horizons
- Robust to moderate model errors
- Choice of regression model is an important design parameter

Continuous on-line planning: iLQR + learned model

Input: base policy π_0

Run base policy to collect data $D \leftarrow \{(s, a, s')_i\}$

Repeat

Fit dynamics model $f(s, a)$ to minimize $\sum_i \|f(s_i, a_i) - s'_i\|^2$

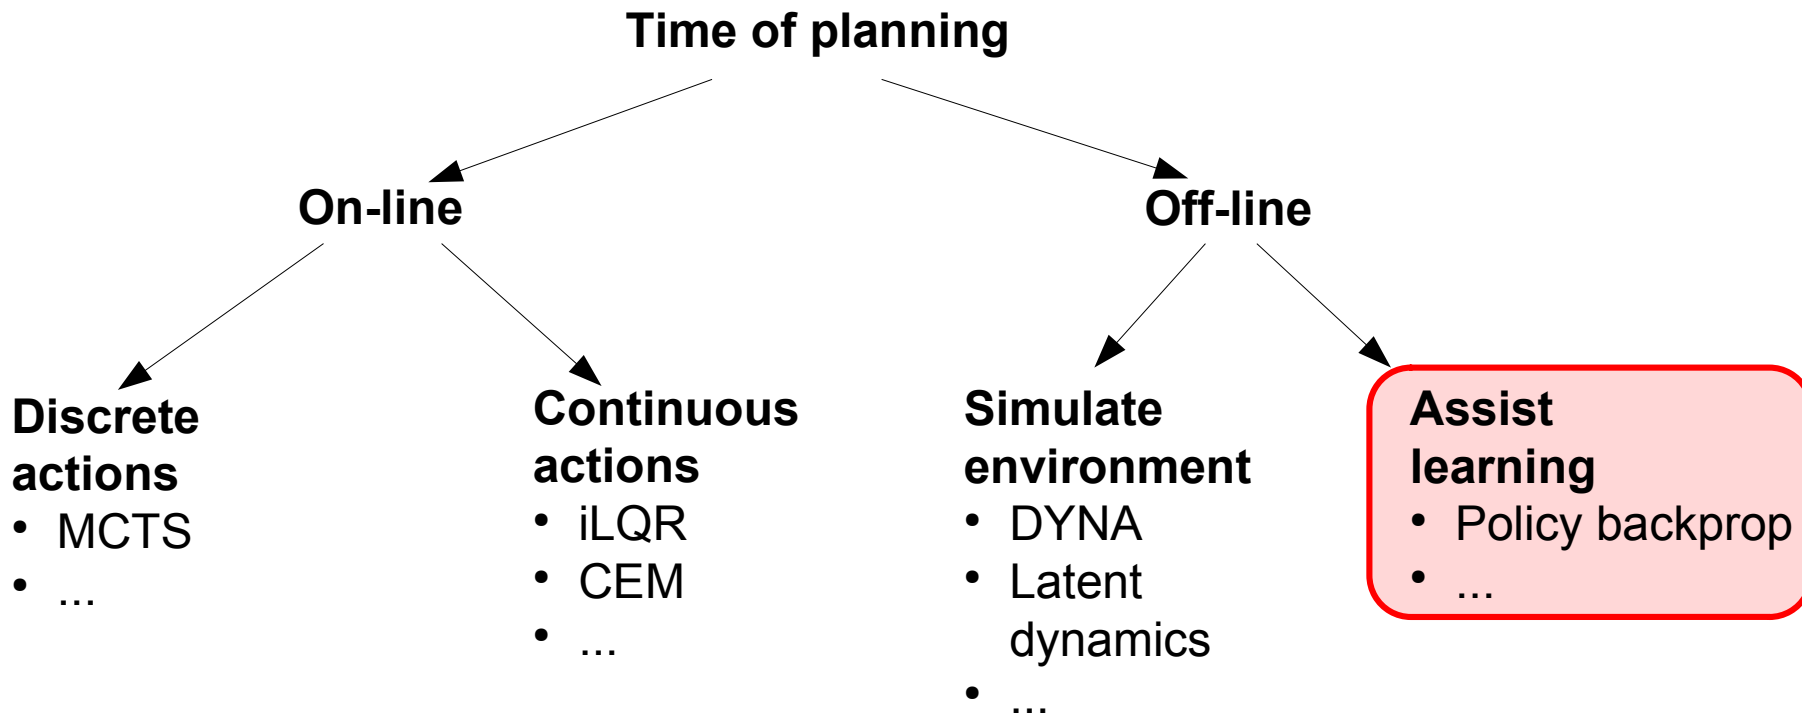
Use model to plan (e.g. iLQR, CEM) actions

Execute first planned action, observe resulting state s'


Update dataset $D \leftarrow D \cup \{(s, a, s')\}$

- Sample efficient.
- Computationally expensive for two reasons
 - Dynamics fitting costly \rightarrow model may be fitted only periodically (every n steps)
 - **Planning costly for long horizons**
- Robust to moderate model errors
- Choice of regression model is an important design parameter

Spectrum of model-based RL



Combining parametric policy with learned dynamics by backpropagation




$$\frac{\partial r_t}{\partial \theta} = \frac{\partial r_t}{\partial a_t} \frac{\partial a_t}{\partial \theta} + \frac{\partial r_t}{\partial s_t} \frac{\partial s_t}{\partial \theta}$$

$$\frac{\partial s_t}{\partial \theta} = \frac{\partial s_t}{\partial s_{t-1}} \frac{\partial s_{t-1}}{\partial \theta} + \frac{\partial s_t}{\partial a_{t-1}} \frac{\partial a_{t-1}}{\partial \theta}$$

policy	reward	dynamics
$\nabla_{\theta} \pi(s_t)$	$\nabla_a r(s_t, a_t)$	$\nabla_s f(s_{t-1}, a_{t-1})$
	$\nabla_s r(s_t, a_t)$	$\nabla_a f(s_{t-1}, a_{t-1})$

Combining parametric policy with learned dynamics by backpropagation



$$\frac{\partial r_t}{\partial \theta} = \frac{\partial r_t}{\partial a_t} \frac{\partial a_t}{\partial \theta} + \frac{\partial r_t}{\partial s_t} \frac{\partial s_t}{\partial \theta}$$

$$\frac{\partial s_t}{\partial \theta} = \frac{\partial s_t}{\partial s_{t-1}} \frac{\partial s_{t-1}}{\partial \theta} + \frac{\partial s_t}{\partial a_{t-1}} \frac{\partial a_{t-1}}{\partial \theta}$$

Run base policy to collect data $D \leftarrow \{(s, a, s')_i\}$

Repeat

Fit dynamics model $f_\phi(s, a)$ to minimize $\sum_i \|f_\phi(s_i, a_i) - s'_i\|^2$

Calculate policy gradient update by backpropagating through dynamics

Execute updated policy (1 or more steps), collect data

Update dataset $D \leftarrow D \cup \{(s, a, s')\}$

Continuous on-line planning: iLQR + learned model

Input: base policy π_0

Run base policy to collect data $D \leftarrow \{(s, a, s')_i\}$

Repeat

Fit dynamics model $f(s, a)$ to minimize $\sum_i \|f(s_i, a_i) - s'_i\|^2$

Use model to plan (e.g. iLQR, CEM) actions

Execute first planned action, observe resulting state s'


Update dataset $D \leftarrow D \cup \{(s, a, s')\}$

- Sample efficient
- Computationally expensive for two reasons
 - Dynamics fitting costly \rightarrow model may be fitted only periodically (every n steps)
 - Planning costly for long horizons
- Robust to moderate model errors
- **Choice of regression model is an important design parameter**

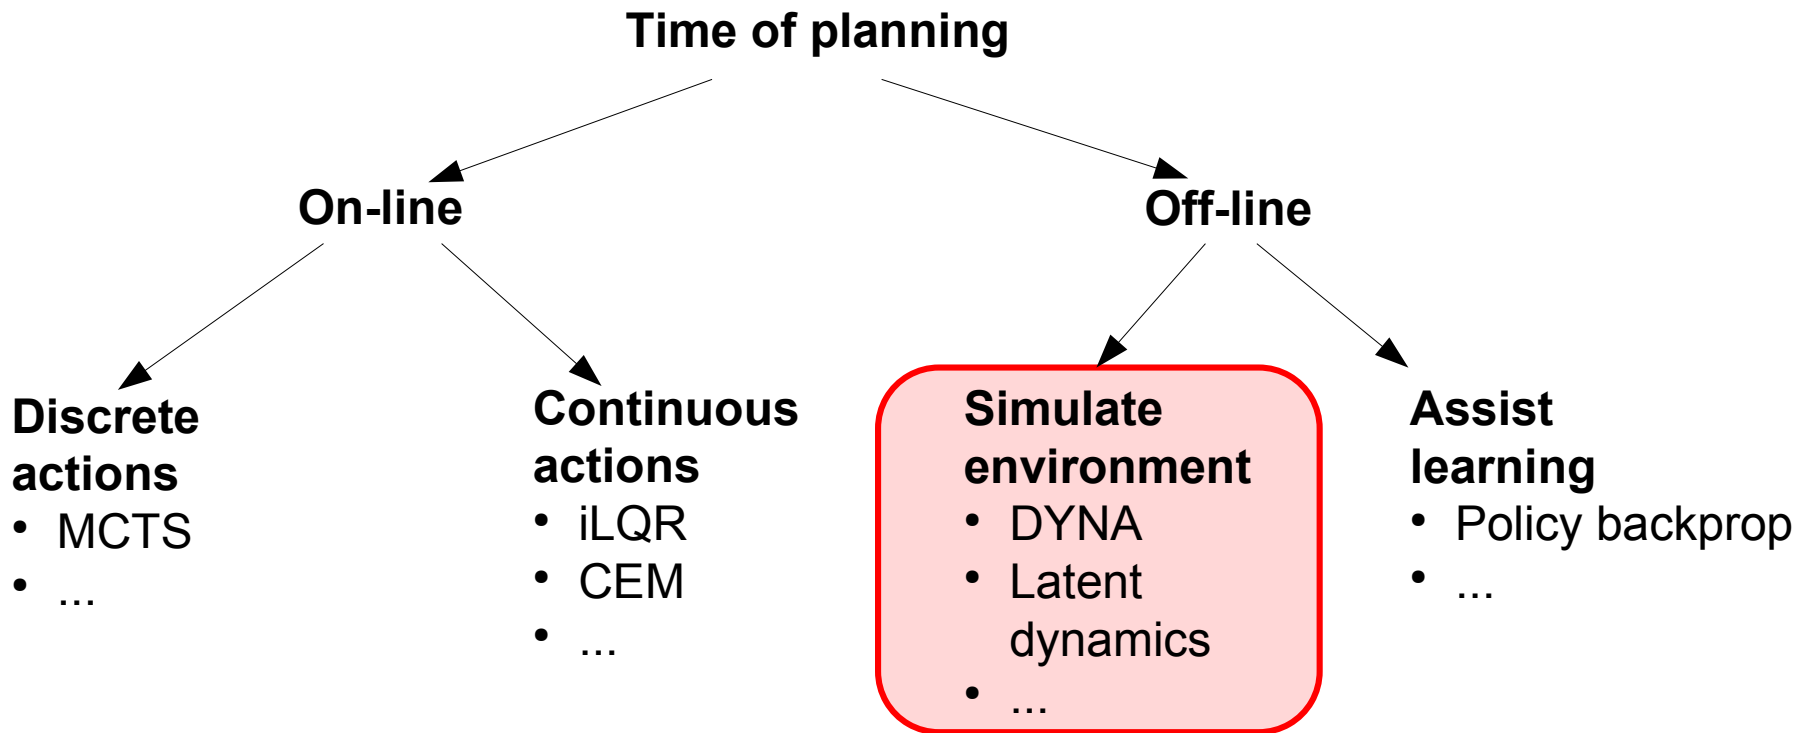
Example

PILCO (Deisenroth&Rasmussen, 2011)

- Dynamics learning: Use Gaussian process models to include model uncertainty. Known quadratic reward
- Simulation: Simulate trajectory with learned model, including uncertainty
- Policy: Radial basis function
- Policy update: Calculate analytically policy gradient using learned dynamics and optimize with quasi-Newton optimizer (BFGS)
- GP → Very sample efficient. Cannot handle a large dataset



Spectrum of model-based RL



Simulate environment to generate additional data: DYNA

Tabular Dyna-Q

Initialize $Q(s, a)$ and $Model(s, a)$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$

Loop forever:

(a) $S \leftarrow$ current (nonterminal) state

(b) $A \leftarrow \varepsilon$ -greedy(S, Q)

(c) Take action A ; observe resultant reward, R , and state, S'

(d) $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

(e) $Model(S, A) \leftarrow R, S'$ (assuming deterministic environment)

(f) Loop repeat n times:

$S \leftarrow$ random previously observed state

$A \leftarrow$ random action previously taken in S

$R, S' \leftarrow Model(S, A)$

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

Update using experience

Update using simulated experience


Learn dynamics model

Generate data by simulating dynamics

Latent dynamics: Motivation

- (real) Dynamics $f(s_{t+1}|s_t, a_t)$
- Reward model $r(r_t|s_t, a_t)$
- Do we need to find an exact dynamics model that is valid for every possible state and action?
- What about learning only a model that allows us to perform the task?
- Some states may share identical optimal policies. Can we take advantage of this somehow?

Learning latent dynamics

- Real dynamics $f(s_{t+1}|s_t, a_t)$
 - Real reward model $r(r_t|s_t, a_t)$
 - Latent state q_t
 - Latent dynamics model $f(q_t|q_{t-1}, a_{t-1}, o_{t-1})$ and $f(q_t|q_{t-1}, a_{t-1})$
 - Latent reward model $r(r_t|q_t)$
 - Policy $\pi(a_t|q_t)$
 - Value function $v(q_t)$
- Observation of the state
- 


Dreamer: learn latent dynamics

- For real world data tuples $(\mathbf{o}_t, \mathbf{a}_t, \mathbf{r}_t)$ update latent state using $f(\mathbf{q}_t | \mathbf{q}_{t-1}, \mathbf{a}_{t-1}, \mathbf{o}_{t-1})$
- and to match real world data update latent models:

$$f(\mathbf{q}_t | \mathbf{q}_{t-1}, \mathbf{a}_{t-1}, \mathbf{o}_{t-1})$$

$$f(\mathbf{q}_t | \mathbf{q}_{t-1}, \mathbf{a}_{t-1})$$


$$r(\mathbf{r}_t | \mathbf{q}_t)$$



Picture adapted from Dream to Control: Learning Behaviors by Latent Imagination [Hafner et al., ICLR 2019]

Dreamer: learn behavior by policy backpropagation


- Simulate latent dynamics using $f(\mathbf{q}_t | \mathbf{q}_{t-1}, \mathbf{a}_{t-1})$
- Estimate value $v(\mathbf{q}_t)$ and rewards
- Update policy $\pi(\mathbf{a}_t | \mathbf{q}_t)$ to maximize value using policy backprop through dynamics (discussed on slide 14)



Picture adapted from Dream to Control: Learning Behaviors by Latent Imagination [Hafner et al., ICLR 2019]

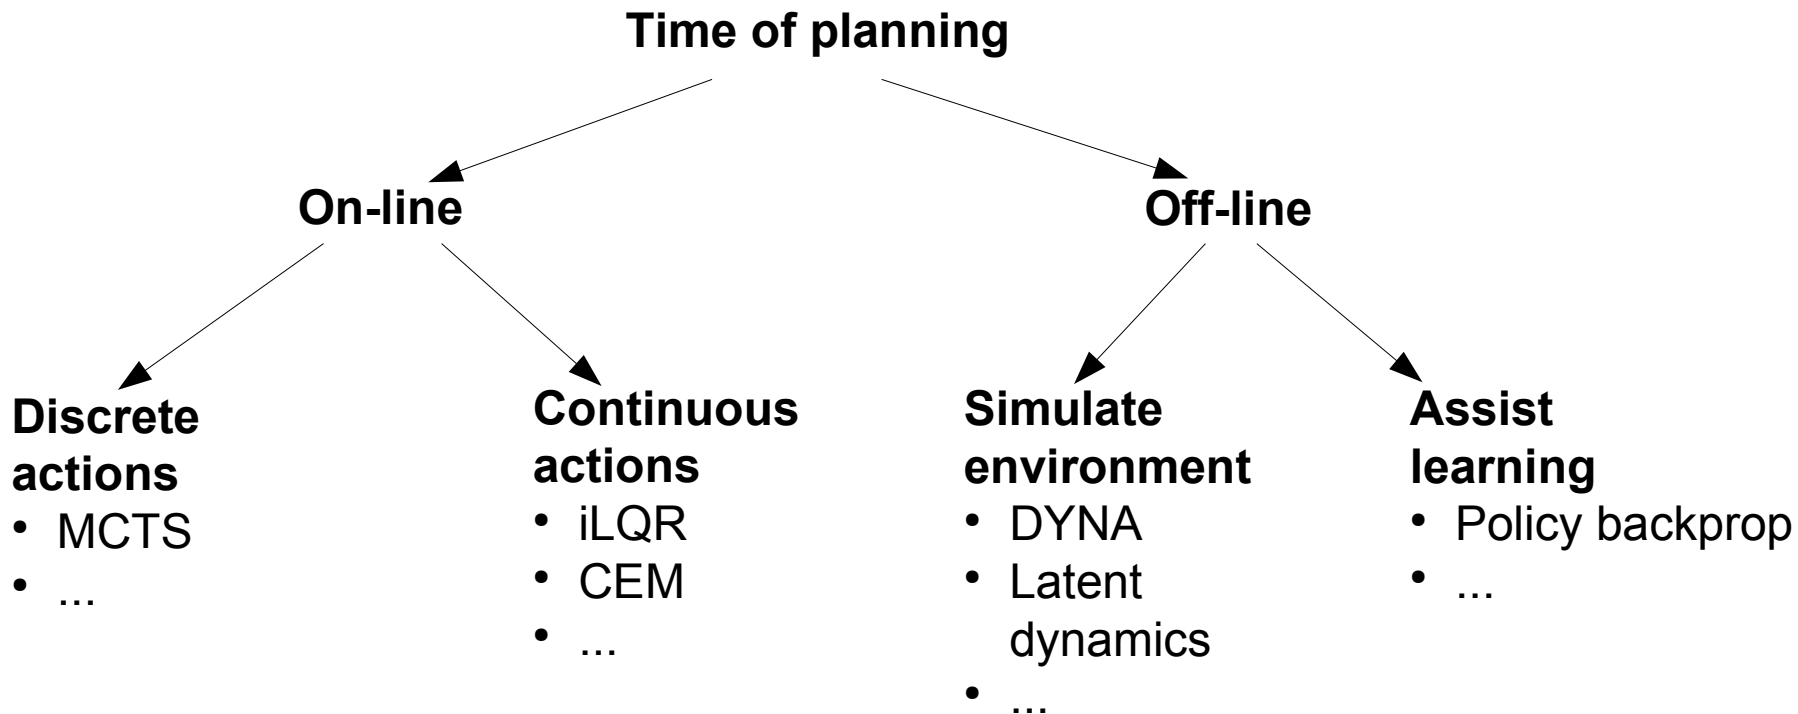
Dreamer: act in the real world

- To collect real world data sample actions from policy $\pi(a_t|q_t)$ and update latent state using $f(q_t|q_{t-1}, a_{t-1}, o_{t-1})$



Picture adapted from Dream to Control: Learning Behaviors by Latent Imagination [Hafner et al., ICLR 2019]

Spectrum of model-based RL



Summary

- Model-based RL requires typically less data than value-based or policy gradient approaches
- Sometimes learned dynamics can be transferred across tasks
- Potentially suboptimal: policy optimization with approximate models may lead to suboptimal solutions and approximate methods to local minima
- Sometimes models are harder to learn than policy
- Often explicit choices required (e.g. time horizon)

Next: exploration / exploitation

- Next week: how to choose actions to find optimal policy?
 - Choose always the best action?
 - But we do not know the best action before we try actions out!
 - How to balance exploration (trying out) with exploitation (choosing what seems the best at the moment)?
 - Monte Carlo tree search (MCTS): balancing exploration vs. exploitation in model-based planning